

▼ HU21CSEN0300296, T.BHAVYA SHREE

```
1 # @title HU21CSEN0300296, T.BHAVYA SHREE
2 from google.colab import drive
3 drive.mount('/content/drive')
```

Mounted at /content/drive

▼ loading the data

```
1 # @title loading the data
2 #loading the data
3 import numpy as np
4 import pandas as pd
5 import os
6 print(os.listdir("/content/drive/MyDrive/D1"))
```

['test_dataset', 'validation_dataset', 'train_dataset', 'vgg_model.h5']

▼ Importing required libraries

```
1 # @title Importing required libraries
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tensorflow.keras.preprocessing import image
5 import os
6 from tensorflow.keras.preprocessing.image import ImageDataGenerator
7 from tensorflow.keras.applications import VGG16
8 from tensorflow.keras.models import Sequential
9 from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
10 from tensorflow.keras.optimizers import Adam
11 from tensorflow.keras.callbacks import EarlyStopping
12 from tensorflow.keras.applications import InceptionV3
13
```

```
1 !pip install python-resize-image
```

Collecting python-resize-image
 Downloading python_resize_image-1.1.20-py2.py3-none-any.whl (8.4 kB)
 Requirement already satisfied: Pillow>=5.1.0 in /usr/local/lib/python3.10/dist-packages (from python-resize-image) (9.4.0)
 Requirement already satisfied: requests>=2.19.1 in /usr/local/lib/python3.10/dist-packages (from python-resize-image) (2.31.0)
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.1->python-resize-image) (3.2.0)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.1->python-resize-image) (3.6)
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.1->python-resize-image) (2.0.7)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.1->python-resize-image) (2024.7.4)
 Installing collected packages: python-resize-image
 Successfully installed python-resize-image-1.1.20

```
1 train_ds = train_datagen.flow_from_directory(
2     train_dataset,
3     class_mode='categorical',
4     shuffle=True,
5     batch_size=32,
6     color_mode='rgb',
7     target_size=(224, 224),
8     subset='training'
9 )
```

Found 18 images belonging to 3 classes.

▼ Exploratory data analysis

```
1 # @title Exploratory data analysis
2 # Checking class distribution in the training dataset
3 class_counts = train_ds.classes
4 unique_classes, class_counts = np.unique(class_counts, return_counts=True)
5 class_distribution = dict(zip(train_ds.class_indices.keys(), class_counts))
6
7 print("Class Distribution:")
8 for cls, count in class_distribution.items():
9     print(f"{cls}: {count} images")
```

Class Distribution:
 drone: 6 images
 flight: 6 images
 helicopter: 6 images

```

1 # Displaying sample images from each class
2 plt.figure(figsize=(15, 10))
3 for i in range(len(train_ds.class_indices)):
4     cls_name = list(train_ds.class_indices.keys())[i]
5     cls_dir = os.path.join(train_dataset, cls_name)
6     sample_img = os.listdir(cls_dir)[0]
7     img_path = os.path.join(cls_dir, sample_img)
8
9     plt.subplot(1, len(train_ds.class_indices), i+1)
10    plt.imshow(image.load_img(img_path))
11    plt.title(cls_name)
12    plt.axis('off')
13 plt.show()

```



drone



flight



helicopter



```

1 # Checking image dimensions
2 sample_img, _ = train_ds[0]
3 img_shape = sample_img[0].shape
4 print(f"Image Shape: {img_shape}")

```



Image Shape: (224, 224, 3)

```

1 # Visualizing augmented images
2 augmented_images = [train_ds[0][0][0] for _ in range(5)]
3
4 plt.figure(figsize=(15, 5))
5 for i, img in enumerate(augmented_images):
6     plt.subplot(1, 5, i+1)
7     plt.imshow(img)
8     plt.axis('off')
9 plt.show()

```



```

1 # path to the test dataset
2 test_dataset = "/content/drive/MyDrive/D1/test_dataset/test_dataset/test_dataset"
3 class_name = list(os.listdir(test_dataset)[0])
4
5 # Initialize the ImageDataGenerator for testing (only rescaling)
6 test_datagen = ImageDataGenerator(rescale=1./255)
7
8 # Load Test Data
9 test_ds = test_datagen.flow_from_directory(
10     test_dataset,
11     class_mode='categorical',
12     shuffle=False,
13     batch_size=1, # Adjust to the actual number of images in your test dataset
14     color_mode='rgb',

```

```

15     target_size=(224, 224) # Adjust the target size for VGG16
16 )
17
18 # Get a batch of images for visualization
19 sample_images, _ = next(test_ds)
20
21 # Visualize a few sample images
22 plt.figure(figsize=(15, 10))
23 for i in range(len(sample_images)):
24     plt.subplot(1, len(sample_images), i + 1)
25     plt.imshow(sample_images[i])
26     plt.title(f"Class: {class_name[i]}")
27     plt.axis('off')
28
29 plt.show()

```

Found 6 images belonging to 3 classes.

Class: d



```

1 from keras.preprocessing import image
2 from keras.preprocessing.image import ImageDataGenerator
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 # Path to your model and sample image
7 model_path = '/content/drive/MyDrive/DL/your_model.h5' # Replace with the path to your model
8 sample_image_path = '/content/drive/MyDrive/DL/test_dataset/test_dataset/ship/677562.jpg' # Replace with the path to your sample image
9
10 # Load your model
11 from keras.models import load_model
12 model = load_model(model_path)
13
14 # Load the sample image
15 sample_image = image.load_img(sample_image_path)
16 img_array = image.img_to_array(sample_image)
17 img_array = np.expand_dims(img_array, axis=0)
18
19 # ImageDataGenerator
20 datagen = ImageDataGenerator(
21     rotation_range=40,

```

```

22     width_shift_range=0.2,
23     height_shift_range=0.2,
24     shear_range=0.2,
25     zoom_range=0.2,
26     horizontal_flip=True,
27     fill_mode='nearest'
28 )
29
30 # Generate augmented images
31 augmented_images = []
32
33 for batch in datagen.flow(img_array, batch_size=1):
34     augmented_images.append(image.array_to_img(batch[0]))
35     if len(augmented_images) >= 5: # Generate 5 augmented images
36         break
37
38 # Display original and augmented images
39 plt.figure(figsize=(27, 4 * (len(augmented_images) + 1)))
40
41 # Original image
42 plt.subplot(len(augmented_images) + 1, 6, 1)
43 plt.imshow(sample_image)
44 plt.title('Original Image')
45
46 # Augmented images on separate lines
47 for i, augmented_img in enumerate(augmented_images):
48     plt.subplot(len(augmented_images) + 1, 6, 6 * (i + 1) + 1)
49     plt.imshow(augmented_img)
50     plt.title(f'Augmented Image {i + 1}')
51
52 # Prediction on original and augmented images
53 plt.subplot(len(augmented_images) + 1, 6, 6)
54 original_prediction = model.predict(img_array)
55 plt.text(0, 0.5, f'Prediction: {original_prediction}', fontsize=12, va='center')
56
57 for i, augmented_img in enumerate(augmented_images):
58     plt.subplot(len(augmented_images) + 1, 6, 6 * (i + 1) + 6)
59     augmented_img_array = image.img_to_array(augmented_img)
60     augmented_img_array = np.expand_dims(augmented_img_array, axis=0)
61     augmented_prediction = model.predict(augmented_img_array)
62     plt.text(0, 0.5, f'Prediction: {augmented_prediction}', fontsize=12, va='center')
63
64 plt.show()
65

```









```

1 from keras.preprocessing import image
2 from keras.preprocessing.image import ImageDataGenerator
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 sample_image_path = '/content/drive/MyDrive/DL/train_dataset/train_dataset/train_dataset/drone/drone 1.png'
7 sample_image = image.load_img(sample_image_path)
8 img_array = image.img_to_array(sample_image)
9 img_array = np.expand_dims(img_array, axis=0)
10
11 #ImageDataGenerator
12 datagen = ImageDataGenerator(
13     rotation_range=40,
14     width_shift_range=0.2,
15     height_shift_range=0.2,
16     shear_range=0.2,
17     zoom_range=0.2,
18     horizontal_flip=True,
19     fill_mode='nearest'
20 )
21
22 # Generate augmented images
23 augmented_images = []
24
25 for batch in datagen.flow(img_array, batch_size=1):
26     augmented_images.append(image.array_to_img(batch[0]))
27     if len(augmented_images) >= 5: # Generate 5 augmented images
28         break
29
30 #original and augmented images
31 plt.figure(figsize=(27, 4 * (len(augmented_images) + 1))) # Adjust the height
32
33 #original image
34 plt.subplot(len(augmented_images) + 1, 6, 1)
35 plt.imshow(sample_image)
36 plt.title('Original Image')
37

```

```
38 # augmented image are then displayed on seperate lines for better visibility
39 for i, augmented_img in enumerate(augmented_images):
40     plt.subplot(len(augmented_images) + 1, 6, 6 * (i + 1) + 1)
41     plt.imshow(augmented_img)
42     plt.title(f'Augmented Image {i + 1}')
43
44 #The augmentation includes random rotations, shifts, shearing, zooming, and horizontal flipping
45 plt.show()
```





```

1 from PIL import Image
2 image_path = '/content/drive/MyDrive/Dl/train_dataset/train_dataset/train_dataset/drone/drone 1.png'
3 pil_image = Image.open(image_path)
4
5 #PIL image to NumPy array
6 image_array = np.array(pil_image)
7
8 #plot of original ship image
9 plt.figure(figsize=(10, 5))
10 plt.subplot(1, 2, 1)
11 plt.imshow(image_array)
12 plt.title('Original Image')
13 plt.axis('off')
14
15 #The second plot with interpolation (has smoothing effect applied by the 'hanning' interpolation)
16 plt.figure(figsize=(5, 4))
17 plt.imshow(image_array, interpolation='hanning')
18 plt.axis('off')
19 plt.show()

```



Original Image



```

1 i, (im1, im2, im3, im4) = plt.subplots(1, 4, sharey=True)
2 i.set_figwidth(20)

```

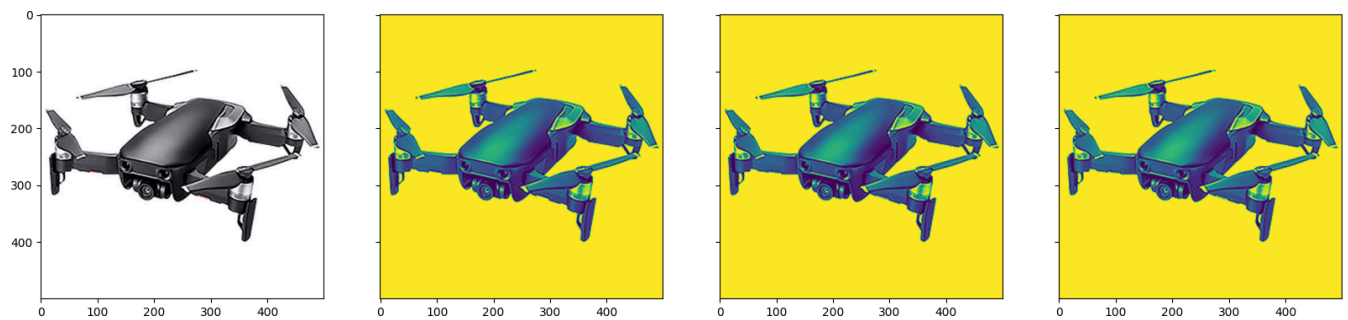
```

3
4 im1.imshow(image_array,interpolation='hanning') #Original image
5 im2.imshow(image_array[:, :, 0],interpolation='hanning') #Red channel
6 im3.imshow(image_array[:, :, 1],interpolation='hanning') #Green channel
7 im4.imshow(image_array[:, :, 2],interpolation='hanning') #Blue channel
8 i.suptitle('Original & RGB image channels')

```

Text(0.5, 0.98, 'Original & RGB image channels')

Original & RGB image channels




```

1 #convert the original RGB image to a grayscale image
2 # used to visualize the intensity (luminance) of an image without considering color information
3 from skimage import io
4 import skimage
5 gray_image = skimage.color.rgb2gray(image_array[:, :, :3])
6 plt.imshow(gray_image, cmap = 'gray', interpolation='hanning')

```

<matplotlib.image.AxesImage at 0x7bb87e0bf9d0>




```

1 norm_image = (gray_image - np.min(gray_image)) / (np.max(gray_image) - np.min(gray_image))
2 plt.imshow(norm_image,interpolation='hanning')


```


 <matplotlib.image.AxesImage at 0x7bb878733820>




```
1 # Import libraries and use Keras' ImageDataGenerator for image data augmentation
2 from numpy import expand_dims
3 from keras.preprocessing.image import load_img
4 from keras.preprocessing.image import ImageDataGenerator
5 import matplotlib.pyplot as plt
6
7 # Expand dimension to one sample
8 samples = expand_dims(image_array, 0)
9
10 # Create image data augmentation generator
11 datagen = ImageDataGenerator(width_shift_range=[5, 10])
12
13 # Create an iterator
14 it = datagen.flow(samples, batch_size=1)
15 fig, im = plt.subplots(nrows=1, ncols=3, figsize=(9, 5))
16
17 # Generate batch of images
18 for i in range(3):
19     # Convert to unsigned integers
20     images = next(it)[0].astype('uint8')
21
22     # Plot image
23     im[i].imshow(images, interpolation='hanning')
24
25 plt.show()
26
```






```
1 datagen = ImageDataGenerator(horizontal_flip=True, vertical_flip=True)
2
3 # create an iterator
4 it = datagen.flow(samples, batch_size=1)
5 fig, im = plt.subplots(nrows=1, ncols=3, figsize=(10,10))
6
7 # generate batch of images
8 for i in range(3):
9
10     # convert to unsigned integers
11     images = next(it)[0].astype('uint8')
12
13     # plot image
14     im[i].imshow(images, interpolation='hanning')
```



```

1 # ImageDataGenerator for brightness
2 datagen = ImageDataGenerator(brightness_range=[0.3,2.0])
3 # create an iterator
4 it = datagen.flow(samples, batch_size=1)
5 fig, im = plt.subplots(nrows=1, ncols=3, figsize=(10,10))
6
7 # generate batch of images
8 for i in range(3):
9
10     # convert to unsigned integers
11     images = next(it)[0].astype('uint8')
12
13     # plot image
14     im[i].imshow(images,interpolation='hanning')

```




```

1 from skimage import io, filters, color, data
2 import os
3 plt.figure(figsize=(5,5))
4 img_sobel = filters.sobel(norm_image)
5 plt.imshow(1-img_sobel, cmap="gray")

```

<matplotlib.image.AxesImage at 0x7bb877b28280>




```

1 plt.figure(figsize=(5,5))
2 img_prewit = filters.prewitt(norm_image)
3 plt.imshow(1-img_prewit, cmap="gray")


```

 <matplotlib.image.AxesImage at 0x7bb8785372e0>




```
1 plt.figure(figsize=(5,5))
2 img_gau = filters.gaussian(image_array,channel_axis=-1)
3 plt.imshow(1-img_gau)
```

 <matplotlib.image.AxesImage at 0x7bb8786d43d0>



```
1 #canny edge detection
2 import cv2 as cv
3 def hist(histing):
4     ycrb = cv.cvtColor(image_array, cv.COLOR_RGB2YCR_CB)
5     channels = cv.split(ycrb)
6     cv.equalizeHist(channels[0], channels[0])
7     cv.merge(channels, ycrb)
8     cv.cvtColor(ycrb, cv.COLOR_YCR_CB2RGB, histing)
9     return histing
10
11 dehizing = hist(norm_image)
12
13 edges = cv.Canny(image_array,50,100)
14
15 plt.subplot(121),plt.imshow(image_array,cmap = 'gray')
16 plt.title('Original Image'), plt.xticks([]), plt.yticks([])
17 plt.subplot(122),plt.imshow(edges,cmap = 'gray')
18 plt.title('Edge Image'), plt.xticks([]), plt.yticks([])
19
20 plt.show()
```



✓ Path for training, testing and validation datasets

```
1 # @title Path for training, testing and validation datasets
2 train_dataset = "/content/drive/MyDrive/DL/train_dataset/train_dataset/train_dataset"
3 test_dataset = "/content/drive/MyDrive/DL/test_dataset/test_dataset/test_dataset"
4 validation_dataset = "/content/drive/MyDrive/DL/validation_dataset/validation_dataset/validation_dataset"
```

✓ Image preprocessing

```
1 # @title Image preprocessing
2 train_datagen = ImageDataGenerator(
3     rescale=1./255,
4     rotation_range=20,
5     horizontal_flip=True,
6     validation_split=0.1
7 )
```

✓ loading data

```
1 # @title loading data
2 # Load Data
3 train_ds = train_datagen.flow_from_directory(
4     train_dataset,
5     class_mode='categorical',
6     shuffle=True,
7     batch_size=32,
8     color_mode='rgb',
9     target_size=(224, 224), # Adjust the target size for VGG16
10    subset='training'
11 )
12
13 valid_ds = train_datagen.flow_from_directory(
14     validation_dataset,
15     class_mode='categorical',
16     shuffle=True,
17     batch_size=32,
18     color_mode='rgb',
19     target_size=(224, 224), # Adjust the target size for VGG16
20    subset='validation'
21 )
```

Found 18 images belonging to 3 classes.
Found 0 images belonging to 3 classes.

✓ VGG16 Model

```
1 # @title VGG16 Model
2 from tensorflow.keras.applications import VGG16
3 base_model = VGG16(include_top=False, input_shape=(224, 224, 3)) # Change the model to VGG16
4 base_model.trainable = False
5
6 # Model Architecture
7 model = Sequential([
8     base_model,
9     GlobalAveragePooling2D(),
10    Dense(256, activation='relu', kernel_initializer='he_normal'),
11    Dropout(0.4),
12    Dense(len(train_ds.class_indices), activation='softmax')
13 ], name="VGG16-TL")
14
15 model.summary()
16
```

```

17 # Compiling Model
18 model.compile(
19     loss='categorical_crossentropy',
20     optimizer=Adam(),
21     metrics=['accuracy']
22 )
23 callbacks = [EarlyStopping(patience=3, restore_best_weights=True)]
24 history = model.fit(train_ds, validation_data=valid_ds, epochs=10, callbacks=callbacks)

```

Model: "VGG16-TL"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 512)	0
dense_12 (Dense)	(None, 256)	131328
dropout_6 (Dropout)	(None, 256)	0
dense_13 (Dense)	(None, 3)	771

=====

Total params: 14846787 (56.64 MB)
 Trainable params: 132099 (516.01 KB)
 Non-trainable params: 14714688 (56.13 MB)

```

Epoch 1/10
1/1 [=====] - ETA: 0s - loss: 1.2802 - accuracy: 0.2778WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 1.2802 - accuracy: 0.2778
Epoch 2/10
1/1 [=====] - ETA: 0s - loss: 1.2561 - accuracy: 0.2778WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 1.2561 - accuracy: 0.2778
Epoch 3/10
1/1 [=====] - ETA: 0s - loss: 1.1706 - accuracy: 0.3889WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 11s 11s/step - loss: 1.1706 - accuracy: 0.3889
Epoch 4/10
1/1 [=====] - ETA: 0s - loss: 1.1109 - accuracy: 0.3889WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 1.1109 - accuracy: 0.3889
Epoch 5/10
1/1 [=====] - ETA: 0s - loss: 0.8140 - accuracy: 0.7222WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 0.8140 - accuracy: 0.7222
Epoch 6/10
1/1 [=====] - ETA: 0s - loss: 1.1160 - accuracy: 0.4444WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 10s 10s/step - loss: 1.1160 - accuracy: 0.4444
Epoch 7/10
1/1 [=====] - ETA: 0s - loss: 0.8825 - accuracy: 0.7778WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 0.8825 - accuracy: 0.7778
Epoch 8/10
1/1 [=====] - ETA: 0s - loss: 0.8168 - accuracy: 0.6667WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 0.8168 - accuracy: 0.6667
Epoch 9/10
1/1 [=====] - ETA: 0s - loss: 0.7871 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 10s 10s/step - loss: 0.7871 - accuracy: 0.5556
Epoch 10/10
1/1 [=====] - ETA: 0s - loss: 0.9508 - accuracy: 0.6111WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 12s 12s/step - loss: 0.9508 - accuracy: 0.6111

```

✓ Testing accuracy and loss of VGG16 Model

```

1 # @title Testing accuracy and loss of VGG16 Model
2 # Load the test dataset
3 test_datagen = ImageDataGenerator(rescale=1./255)
4
5 test_ds = test_datagen.flow_from_directory(
6     test_dataset,
7     class_mode='categorical',
8     shuffle=False,
9     batch_size=32,
10    color_mode='rgb',
11    target_size=(224, 224) # Adjust the target size for VGG16
12 )
13
14 # Evaluate the model on the test dataset
15 test_loss, test_accuracy = model.evaluate(test_ds)
16
17 print(f'Test Loss: {test_loss:.4f}')
18 print(f'Test Accuracy: {test_accuracy * 100:.2f}%')

```

Found 6 images belonging to 3 classes.

```

1/1 [=====] - 3s 3s/step - loss: 0.8455 - accuracy: 1.0000
Test Loss: 0.8455

```

Test Accuracy: 100.00%

Testing the VGG16 Model

```

1 # @title Testing the VGG16 Model
2 # Set the path to your test dataset
3 test_dataset = "/content/drive/MyDrive/DL/test_dataset/test_dataset/test_dataset"
4
5 # Initialize the ImageDataGenerator for testing (only rescaling)
6 test_datagen = ImageDataGenerator(rescale=1./255)
7
8 # Load Test Data
9 test_ds = test_datagen.flow_from_directory(
10     test_dataset,
11     class_mode='categorical',
12     shuffle=False,
13     batch_size=1, # Adjust to the actual number of images in your test dataset
14     color_mode='rgb',
15     target_size=(224, 224) # Adjust the target size for VGG16
16 )
17
18 # Get class indices
19 class_indices = {v: k for k, v in test_ds.class_indices.items()}
20
21 # Visualize predictions on a few test images
22 plt.figure(figsize=(15, 10))
23 for i in range(len(test_ds)):
24     plt.subplot(4, 5, i+1)
25     img, true_label = test_ds[i]
26     pred_probs = model.predict(img)
27     pred_label = np.argmax(pred_probs, axis=1)[0]
28
29     plt.imshow(img[0])
30     plt.title(f"True: {class_indices[true_label[0].argmax()]} \n Pred: {class_indices[pred_label]}")
31     plt.axis('off')
32
33 plt.tight_layout()
34 plt.show()

```

Found 6 images belonging to 3 classes.

```

1/1 [=====] - 1s 933ms/step
1/1 [=====] - 0s 477ms/step
1/1 [=====] - 0s 492ms/step
1/1 [=====] - 0s 481ms/step
1/1 [=====] - 0s 491ms/step
1/1 [=====] - 0s 495ms/step

```



```

1 from tensorflow.keras.applications import InceptionV3
2 # Set the paths to your dataset
3 train_dataset = "/content/drive/MyDrive/DL/train_dataset/train_dataset/train_dataset"
4 test_dataset = "/content/drive/MyDrive/DL/test_dataset/test_dataset/test_dataset"
5 validation_dataset = "/content/drive/MyDrive/DL/validation_dataset/validation_dataset/validation_dataset"

```

InceptionV3 Model

```

1 # @title InceptionV3 Model
2 # Load Pretrained InceptionV3 Model
3 base_model = InceptionV3(include_top=False, input_shape=(256, 256, 3))
4 base_model.trainable = False
5
6 # Model Architecture
7 model = Sequential([
8     base_model,
9     GlobalAveragePooling2D(),
10    Dense(256, activation='relu', kernel_initializer='he_normal'),
11    Dropout(0.4),
12    Dense(len(train_ds.class_indices), activation='softmax')
13 ], name="Inception-TL")
14
15 model.summary()
16
17 # Compile Model
18 model.compile(
19     loss='categorical_crossentropy',
20     optimizer=Adam(),
21     metrics=['accuracy']
22 )
23
24 # Callbacks
25 callbacks = [EarlyStopping(patience=3, restore_best_weights=True)]
26
27 # Model Training
28 history = model.fit(train_ds, validation_data=valid_ds, epochs=10, callbacks=callbacks)

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception_v3/inception_v3_weights_tf_dim_ordering_87910968/87910968 [=====] - 1s 0us/step
Model: "Inception-TL"

Layer (type)	Output Shape	Param #
=====		
inception_v3 (Functional)	(None, 6, 6, 2048)	21802784
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 2048)	0
dense_2 (Dense)	(None, 256)	524544
dropout_1 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 3)	771
=====		
Total params: 22328099 (85.17 MB)		
Trainable params: 525315 (2.00 MB)		
Non-trainable params: 21802784 (83.17 MB)		

```

Epoch 1/10
1/1 [=====] - ETA: 0s - loss: 1.4728 - accuracy: 0.2778WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 14s 14s/step - loss: 1.4728 - accuracy: 0.2778
Epoch 2/10
1/1 [=====] - ETA: 0s - loss: 0.6912 - accuracy: 0.7222WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 4s 4s/step - loss: 0.6912 - accuracy: 0.7222
Epoch 3/10
1/1 [=====] - ETA: 0s - loss: 0.3377 - accuracy: 0.8889WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 3s 3s/step - loss: 0.3377 - accuracy: 0.8889
Epoch 4/10
1/1 [=====] - ETA: 0s - loss: 0.2638 - accuracy: 0.8889WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 4s 4s/step - loss: 0.2638 - accuracy: 0.8889
Epoch 5/10
1/1 [=====] - ETA: 0s - loss: 0.1216 - accuracy: 0.9444WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 4s 4s/step - loss: 0.1216 - accuracy: 0.9444
Epoch 6/10
1/1 [=====] - ETA: 0s - loss: 0.1675 - accuracy: 0.9444WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 3s 3s/step - loss: 0.1675 - accuracy: 0.9444
Epoch 7/10
1/1 [=====] - ETA: 0s - loss: 0.0291 - accuracy: 1.0000WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 4s 4s/step - loss: 0.0291 - accuracy: 1.0000
Epoch 8/10
1/1 [=====] - ETA: 0s - loss: 0.0317 - accuracy: 1.0000WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 5s 5s/step - loss: 0.0317 - accuracy: 1.0000
Epoch 9/10
1/1 [=====] - ETA: 0s - loss: 0.0120 - accuracy: 1.0000WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 3s 3s/step - loss: 0.0120 - accuracy: 1.0000
Epoch 10/10
1/1 [=====] - ETA: 0s - loss: 0.0149 - accuracy: 1.0000WARNING:tensorflow:Early stopping conditioned on metric
1/1 [=====] - 3s 3s/step - loss: 0.0149 - accuracy: 1.0000

```

✓ Testing the InceptionV3 Model

```

1 # @title Testing the InceptionV3 Model
2 # Set the path to your test dataset
3 test_dataset = "/content/drive/MyDrive/Dl/test_dataset/test_dataset/test_dataset"
4
5 # Initialize the ImageDataGenerator for testing (only rescaling)
6 test_datagen = ImageDataGenerator(rescale=1./255)
7
8 # Load Test Data
9 test_ds = test_datagen.flow_from_directory(
10     test_dataset,
11     class_mode='categorical',
12     shuffle=False,
13     batch_size=1, # Adjust to the actual number of images in your test dataset
14     color_mode='rgb',
15     target_size=(256, 256)
16 )
17
18 # Evaluate the model on the test set
19 test_loss, test_accuracy = model.evaluate(test_ds)
20 print(f'InceptionV3 Test Loss: {inception_test_loss:.4f}')
21 print(f'InceptionV3 Test Accuracy: {inception_test_accuracy * 100:.2f}%')
22 # Get class indices
23 class_indices = {v: k for k, v in test_ds.class_indices.items()}
24
25 # Visualize predictions on a few test images
26 plt.figure(figsize=(15, 10))
27 for i in range(len(test_ds)):
28     plt.subplot(4, 5, i+1)
29     img, true_label = test_ds[i]
30     pred_probs = model.predict(img)
31     pred_label = np.argmax(pred_probs, axis=1)[0]
32
33     plt.imshow(img[0])
34     plt.title(f"True: {class_indices[true_label[0].argmax()]} \n Pred: {class_indices[pred_label]}")
35     plt.axis('off')
36
37 plt.tight_layout()
38 plt.show()

```

Found 6 images belonging to 3 classes.

6/6 [=====] - 5s 755ms/step - loss: 0.8641 - accuracy: 1.0000

InceptionV3 Test Loss: 0.8455

InceptionV3 Test Accuracy: 100.00%

1/1 [=====] - 1s 782ms/step

1/1 [=====] - 1s 640ms/step

1/1 [=====] - 1s 623ms/step

1/1 [=====] - 1s 632ms/step

1/1 [=====] - 1s 638ms/step

1/1 [=====] - 1s 622ms/step

True: drone
Pred: drone



True: drone
Pred: drone



True: flight
Pred: flight



True: flight
Pred: flight



True: helicopter
Pred: helicopter



True: helicopter
Pred: helicopter



▼ Comparing the VGG16 and InceptionV3

```

1 # @title Comparing the VGG16 and InceptionV3
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.metrics import confusion_matrix, classification_report
5 # Evaluate the VGG16 model on the test set

```



```

6 vgg_test_loss, vgg_test_accuracy = model.evaluate(test_ds)
7 print(f'VGG16 Test Loss: {vgg_test_loss:.4f}')
8 print(f'VGG16 Test Accuracy: {vgg_test_accuracy * 100:.2f}%')
9
10 # Evaluate the InceptionV3 model on the test set
11 inception_test_loss, inception_test_accuracy = model.evaluate(test_ds)
12 print(f'InceptionV3 Test Loss: {inception_test_loss:.4f}')
13 print(f'InceptionV3 Test Accuracy: {inception_test_accuracy * 100:.2f}%')

6/6 [=====] - 4s 614ms/step - loss: 0.8455 - accuracy: 1.0000
VGG16 Test Loss: 0.8455
VGG16 Test Accuracy: 100.00%
6/6 [=====] - 5s 837ms/step - loss: 0.8455 - accuracy: 1.0000
InceptionV3 Test Loss: 0.8455
InceptionV3 Test Accuracy: 100.00%

1 # VGG16 Model
2
3 # Load Pretrained VGG16 Model
4 base_model_vgg16 = VGG16(include_top=False, input_shape=(224, 224, 3)) # Change the model to VGG16
5 base_model_vgg16.trainable = False
6
7 # Model Architecture
8 model_vgg16 = Sequential([
9     base_model_vgg16,
10    GlobalAveragePooling2D(),
11    Dense(256, activation='relu', kernel_initializer='he_normal'),
12    Dropout(0.4),
13    Dense(len(train_ds.class_indices), activation='softmax')
14 ], name="VGG16-TL")
15
16 model_vgg16.summary()
17
18 # Compile Model
19 model_vgg16.compile(
20     loss='categorical_crossentropy',
21     optimizer=Adam(),
22     metrics=['accuracy']
23 )
24
25 # Training for VGG16
26 history_vgg16 = model_vgg16.fit(train_ds, validation_data=valid_ds, epochs=10, callbacks=callbacks)
27
28 # InceptionV3 Model
29
30 # Load Pretrained InceptionV3 Model
31 base_model_inception = InceptionV3(include_top=False, input_shape=(256, 256, 3))
32 base_model_inception.trainable = False
33
34 # Model Architecture
35 model_inception = Sequential([
36     base_model_inception,
37     GlobalAveragePooling2D(),
38     Dense(256, activation='relu', kernel_initializer='he_normal'),
39     Dropout(0.4),
40     Dense(len(train_ds.class_indices), activation='softmax')
41 ], name="Inception-TL")
42
43 model_inception.summary()
44
45 # Compile Model
46 model_inception.compile(
47     loss='categorical_crossentropy',
48     optimizer=Adam(),
49     metrics=['accuracy']
50 )
51
52 # Training for InceptionV3
53 history_inception = model_inception.fit(train_ds, validation_data=valid_ds, epochs=10, callbacks=callbacks)
54
55 # Evaluation
56
57 # Predict and evaluate on the training set for VGG16
58 train_accuracy_vgg16 = model_vgg16.evaluate(train_ds)[1]
59 print(f"Train Accuracy (VGG16): {train_accuracy_vgg16}")
60
61 # Predict and evaluate on the test set for VGG16
62 test_accuracy_vgg16 = model_vgg16.evaluate(test_ds)[1]
63 print(f"Test Accuracy (VGG16): {test_accuracy_vgg16}")
64
65 # Predict and evaluate on the training set for InceptionV3
66 train_accuracy_inception = model_inception.evaluate(train_ds)[1]
67 print(f"Train Accuracy (InceptionV3): {train_accuracy_inception}")

```

```

68
69 # Predict and evaluate on the test set for InceptionV3
70 test_accuracy_inception = model_inception.evaluate(test_ds)[1]
71 print(f"Test Accuracy (InceptionV3): {test_accuracy_inception}")
72

```

Model: "VGG16-TL"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
global_average_pooling2d_9 (GlobalAveragePooling2D)	(None, 512)	0
dense_18 (Dense)	(None, 256)	131328
dropout_9 (Dropout)	(None, 256)	0
dense_19 (Dense)	(None, 3)	771

```

=====
Total params: 14846787 (56.64 MB)
Trainable params: 132099 (516.01 KB)
Non-trainable params: 14714688 (56.13 MB)

```

```

Epoch 1/10
1/1 [=====] - ETA: 0s - loss: 1.1013 - accuracy: 0.2778WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 15s 15s/step - loss: 1.1013 - accuracy: 0.2778
Epoch 2/10
1/1 [=====] - ETA: 0s - loss: 1.1338 - accuracy: 0.3889WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 10s 10s/step - loss: 1.1338 - accuracy: 0.3889
Epoch 3/10
1/1 [=====] - ETA: 0s - loss: 0.8591 - accuracy: 0.6667WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 12s 12s/step - loss: 0.8591 - accuracy: 0.6667
Epoch 4/10
1/1 [=====] - ETA: 0s - loss: 0.9937 - accuracy: 0.3333WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 10s 10s/step - loss: 0.9937 - accuracy: 0.3333
Epoch 5/10
1/1 [=====] - ETA: 0s - loss: 0.9740 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 12s 12s/step - loss: 0.9740 - accuracy: 0.5556
Epoch 6/10
1/1 [=====] - ETA: 0s - loss: 0.7966 - accuracy: 0.6111WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 10s 10s/step - loss: 0.7966 - accuracy: 0.6111
Epoch 7/10
1/1 [=====] - ETA: 0s - loss: 0.8530 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 12s 12s/step - loss: 0.8530 - accuracy: 0.5556
Epoch 8/10
1/1 [=====] - ETA: 0s - loss: 0.9500 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 11s 11s/step - loss: 0.9500 - accuracy: 0.5556
Epoch 9/10
1/1 [=====] - ETA: 0s - loss: 0.9261 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 12s 12s/step - loss: 0.9261 - accuracy: 0.5556
Epoch 10/10
1/1 [=====] - ETA: 0s - loss: 0.9792 - accuracy: 0.5556WARNING:tensorflow:Early stopping conditioned on
1/1 [=====] - 11s 11s/step - loss: 0.9792 - accuracy: 0.5556
Model: "Inception-TL"

```

Layer (type)	Output Shape	Param #
inception_v3 (Functional)	(None, 6, 6, 2048)	21802784
global average pooling2d_1	(None, 2048)	0

```

1 # Classification metrics and confusion matrix
2 from sklearn.metrics import classification_report, confusion_matrix
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import numpy as np
6
7 # Function to plot training history
8 def plot_training_history(history, model_name):
9     plt.figure(figsize=(12, 4))
10
11     # Plot training & validation accuracy values
12     plt.subplot(1, 2, 1)
13     plt.plot(history.history['accuracy'])
14     plt.plot(history.history['val_accuracy'])
15     plt.title(f'{model_name} Model Accuracy')
16     plt.xlabel('Epoch')
17     plt.ylabel('Accuracy')
18     plt.legend(['Train', 'Validation'], loc='upper left')
19
20     # Plot training & validation loss values
21     plt.subplot(1, 2, 2)

```